

Planning Methods for Near-Optimal Nonlinear Control

Lucian Buşoniu

Automation Department
Technical University of Cluj-Napoca, Romania
Contact: lucian@busoniu.net

30 June 2016, City University London





Overall theme


AI-based control of complex systems

Complexity: nonlinearity, stochastic dynamics, unknown behavior, distributed structure, ...

Applications: robotics, control, medicine, ...




Optimal control problem (deterministic MDP)



- Controller measures **states x** , applies **actions u**
- System: **dynamics $x_{k+1} = f(x_k, u_k)$**
- Performance: **reward function $r_{k+1} = \rho(x_k, u_k)$**
- **Objective**: maximize discounted return $\sum_{k=0}^{\infty} \gamma^k r_{k+1}$, discount factor $\gamma \in (0, 1)$

AI perspective



- Agent observes **state**, applies **action**
- Environment changes state according to dynamics
... and sends back a **reward**, according to reward function
- **Objective:** maximize discounted return

Example: Quanser pendulum



System:

- x = rod angle α , base angle θ , angular velocities
- u = motor voltage $\in [-9, 9] \text{ V}$
- Sampling time $T_s = 0.05$

Goal: stabilize pointing up:

- $\rho = -\alpha^2 - \theta^2 - .005(\dot{\alpha}^2 + \dot{\theta}^2) - .05u^2$, normalized to $[0, 1]$
- Discount factor $\gamma = 0.85$
- Swingup required

Solution methods

By path to solution:

- **Policy iteration**: find the value function of current control policy, use it to improve the policy, repeat to convergence
- **Value iteration**: find optimal value function, use it to compute optimal policy
- **Policy search**: look directly for the optimal policy

By model usage:

- **Model-based**: dynamics and reward function known
- **Model-free**: only data (**reinforcement learning**)

By interaction level:


- **Offline**: algorithm is run in advance
- **Online**: algorithm runs while controlling system



Online planning

At each step k , solve local optimal control at state x_k :

- Infinite action sequences: $\mathbf{u}_\infty = (u_k, u_{k+1}, \dots)$
 - Optimization problem: $\sup_{\mathbf{u}_\infty} v(\mathbf{u}_\infty) (= \sum_{i=0}^{\infty} \gamma^i r_{k+1+i})$
1. Explore sequences from x_k , to find a near-optimal one
 2. Apply first action of this sequence, and repeat



Receding-horizon model-predictive control

Optimistic planning (OP): Main idea

initialize **set of all possible sequences**


repeat

 select most promising, **optimistic** set

 refine selected set

until computation budget n exhausted

return sequence in best set



Bandit-based optimization; branch & bound if deterministic

Advantages of OP

- **Near-optimality guarantees** as a function of computation n and of complexity m of the problem:

$$\text{error} = O(g(n, m))$$

(Munos, 2014)

- ...for general nonlinear dynamics and rewards



- 1 Introduction
- 2 OPC: Optimistic planning with continuous actions
 - Setting and algorithm
 - Analysis
 - Experiments
- 3 SOPC: Simultaneous OPC
 - Algorithm and guarantee
 - Experiments
- 4 Conclusions

Assumptions

- Rewards $r \in [0, 1]$
- Action space $U = [0, 1]$
(can be extended to compact multidimensional U)
- Lipschitz dynamics and rewards:



$$\|f(x, u) - f(x', u')\| \leq L_f(\|x - x'\| + |u - u'|)$$

$$|\rho(x, u) - \rho(x', u')| \leq L_\rho(\|x - x'\| + |u - u'|)$$

- $\gamma L_f < 1$: most restrictive

Search refinement

- Split U^∞ iteratively, leading to a tree of hyperboxes




- Each box i only represents explicitly dimensions already split, $k = 0, \dots, K_i - 1$
- Box i has value $v(i) = \sum_{k=0}^{K_i-1} \gamma^k r_{i,k+1}$, rewards of center sequence

Lipschitz value function

- For any two **action sequences** $\mathbf{u}_\infty, \mathbf{u}'_\infty$:

$$|v(\mathbf{u}_\infty) - v(\mathbf{u}'_\infty)| \leq \frac{L_\rho}{1 - \gamma L_f} \sum_{k=0}^{\infty} \gamma^k |u_k - u'_k|$$

- Intuition: **states** (and so **rewards**) may diverge somewhat, but divergence controlled due to $\gamma L_f < 1$




Box upper bound

- For any sequence \mathbf{u}_∞ in box i :

$$v(\mathbf{u}_\infty) \leq v(i) + \frac{\max\{1, L_\rho\}}{1 - \gamma L_f} \sum_{k=0}^{\infty} \gamma^k d_{i,k} := b(i)$$

- $d_{i,k}$ length of dimension k , 1 if not split yet



- $b(i)$ **b-value** of box i

Diameter and dimension selection

- **Diameter** $\delta(i) := \frac{\max\{1, L_\rho\}}{1 - \gamma L_f} \sum_{k=0}^{\infty} \gamma^k d_{i,k}$
= uncertainty on values in the box
 - **Impact** of dimension k on uncertainty is $\gamma^k d_{i,k}$
- ⇒ when splitting a box, choose dimension with largest impact, to reduce uncertainty the most
- Always split into odd $M > 1/\gamma$ pieces

OPC algorithm

initialize tree with root box U^∞

while n not exhausted **do**

 select **optimistic** leaf box $i^\dagger = \arg \max_{i \in \mathcal{L}} b(i)$

 select **max-impact** dimension $k^\dagger = \arg \max_k \gamma^k d_{i^\dagger, k}$

 split i^\dagger along k^\dagger , creating M children on the tree

end while

return best center sequence seen, $i^* = \arg \max_i v(i)$

(ACC 2016)



- 1 Introduction
- 2 **OPC: Optimistic planning with continuous actions**
 - Setting and algorithm
 - **Analysis**
 - Experiments
- 3 SOPC: Simultaneous OPC
- 4 Conclusions


Diameter bound

Lemma

Given depth in tree h = total number of splits:

$$\delta(i) = \tilde{O}(\gamma \sqrt{2h^{\frac{\tau-1}{\tau^2}}}), \text{ where } \tau = \left\lceil \frac{\log 1/M}{\log \gamma} \right\rceil$$

Diameters vary by the order of splits, but they all converge to 0 roughly exponentially in \sqrt{h} . Example:



Complexity measure


- OPC only expands in near-optimal subtree:

$$\mathcal{T}^* = \{i \in \mathcal{T} \mid v^* - v(i) \leq \delta(i)\}$$

(nodes that cannot be eliminated as suboptimal)

- Define $m \in [1, M]$ = asymptotic branching factor of \mathcal{T}^* :
problem complexity measure

E.g. $m = 2, M = 3$:



Performance guarantee

Theorem

After spending n model calls, OPC suboptimality is:

$$v^* - v(i^*) = \begin{cases} \tilde{O}(\gamma \sqrt{\frac{2(\tau-1) \log n}{\tau^2 \log m}}), & \text{if } m > 1 \\ \tilde{O}(\gamma n^{1/4} b), & \text{if } m = 1 \end{cases}$$

- Convergence rate modulated by problem complexity m , faster when m smaller
- When $m = 1$, convergence is fast, with power $n^{1/4}$
- When $m > 1$, we pay for generality: exponential computation m^h to reach depth h

Related work

- Inspired by **optimistic optimization (DOO & SOO)**

(Munos 2011)

- **HOLOP, HOOT**: fixed finite horizon

(Weinstein et al. 2012, Mansley et al. 2011)

- **Lipschitz planning (LP)**: different dimension selection, no guarantees

(Hren 2012)

- **SOOP**: no Lipschitz constants, no guarantees

(ADPRL 2013)



- 1 Introduction
- 2 **OPC: Optimistic planning with continuous actions**
 - Setting and algorithm
 - Analysis
 - **Experiments**
- 3 SOPC: Simultaneous OPC
- 4 Conclusions

Recall example: Quanser pendulum



System:


- x = rod angle α , base angle θ , angular velocities
- u = motor voltage $\in [-9, 9] \text{ V}$
- Sampling time $T_s = 0.05$

Goal: stabilize pointing up:

- $\rho = -\alpha^2 - \theta^2 - .005(\dot{\alpha}^2 + \dot{\theta}^2) - .05u^2$, normalized to $[0, 1]$
- Discount factor $\gamma = 0.85$
- Swingup required

OPC versus LP and SOOP

OPC with $L_f = L_\rho = 1.1$, tighter diameter formula
Parameters of algorithms optimized





⇒ OPC theoretically nice, beaten by heuristic SOOP in practice
Disadvantage: **under/overestimated Lipschitz constants**

- 1 Introduction
- 2 OPC: Optimistic planning with continuous actions
 - Setting and algorithm
 - Analysis
 - Experiments
- 3 **SOPC: Simultaneous OPC**
 - **Algorithm and guarantee**
 - **Experiments**
- 4 Conclusions

Idea

- Avoid using Lipschitz constants (i.e. diameters) altogether
- ⇒ Split a **potentially optimistic** box at each depth:

$$i_h^\dagger = \arg \max_{i \text{ at } h} v(i); \text{ proxy for unknown } b(i) = v(i) + \delta(i)$$



- Depth cutoff at $h_{\max}(n)$ to avoid indefinite expansion

SOPC algorithm

initialize tree with root box

loop

for $h = \text{first unexpanded to } h_{\max}(n)$ **do**

potentially optimistic leaf $i_h^\dagger = \arg \max_{i \in \mathcal{L}_h} v(i)$

max-impact dimension $k_h^\dagger = \arg \max_k \gamma^k d_{i_h^\dagger, k}$

split i_h^\dagger along k_h^\dagger

if budget n reached, **stop algo** **end if**

end for

end loop

return best sequence seen $i^* = \arg \max_i v(i)$

Performance guarantee

Theorem

For budget n , SOPC suboptimality is:

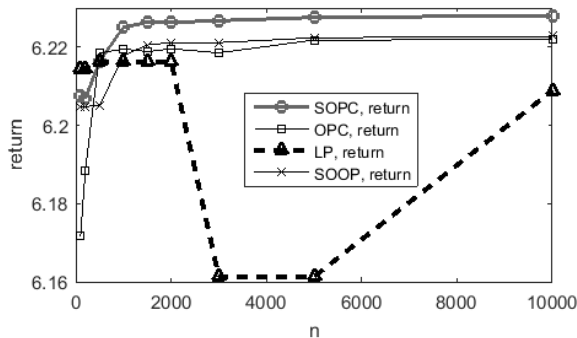
$$v^* - v(i^*) = \begin{cases} \tilde{O}\left(\gamma \sqrt{\frac{2(1-2\varepsilon)(\tau-1) \log n}{\tau^2 \log m}}\right), & \text{if } m > 1 \text{ and } h_{\max}(n) = n^\varepsilon \\ \tilde{O}(\gamma n^{1/6} b), & \text{if } m = 1 \text{ and } h_{\max}(n) = n^{1/3} \end{cases}$$

- When $m > 1$, with small ε nearly same bound as OPC
- Intuition: expanding full path takes up to $h_{\max}(n)$, negligible compared to exponential tree size m^h
- When $m = 1$, $n^{1/6}$ instead of $n^{1/4}$ – slower but similar
- All this while **adapting to unknown smoothness**

- 1 Introduction
- 2 OPC: Optimistic planning with continuous actions
- 3 **SOPC: Simultaneous OPC**
 - Algorithm and guarantee
 - **Experiments**
- 4 Conclusions

Quanser pendulum results


SOPC with $h_{\max}(n) = n^{0.45}$



⇒ best algorithm

Controlled trajectory

$n = 5000$ model calls; note adaptive discretization of control magnitude



Real-time control



Conclusions

(Simultaneous) optimistic planning with continuous actions:

- Control of general nonlinear systems, guaranteed near-optimal
- SOPC adapts to unknown smoothness, works well in practice

Next steps:

- Eliminate assumption $\gamma L_f < 1$ using stability
- Reduce complexity while still keeping problem interesting?

Thank you!

